

# Simplicity Lies in the Eye of the Beholder

---

## A Strategic Perspective on Controllers in Reactive Synthesis

Mickael Randour

F.R.S.-FNRS & UMONS – Université de Mons, Belgium

October 3, 2025

*Reachability Problems 2025*



Special thanks to the *Delegación General Valonia-Bruselas en España*.

# The talk in one slide

# The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

# The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

**Simpler is better:**

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

# The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

**Simpler is better:**

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

# The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

**Simpler is better:**

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

**But how to define complexity and how to measure it?**

# The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

**Simpler is better:**

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

**But how to define complexity and how to measure it?**

↪ **That is our topic of the today.**

# The talk in one **more** slide

## The talk in one **more** slide

**Yes, I lied, and I will lie even more.** The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems. . .

## The talk in one **more** slide

**Yes, I lied, and I will lie even more.** The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

**There will be some hand-waving and approximations to keep the talk high level.**

→ Check the paper for more details.

## The talk in one **more** slide

**Yes, I lied, and I will lie even more.** The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

**There will be some hand-waving and approximations to keep the talk high level.**

→ Check the paper for more details.

→ I will focus on recent work with marvelous co-authors.

## 1 Controller synthesis

## 2 Memory

## 3 Randomness

## 4 Beyond Mealy machines

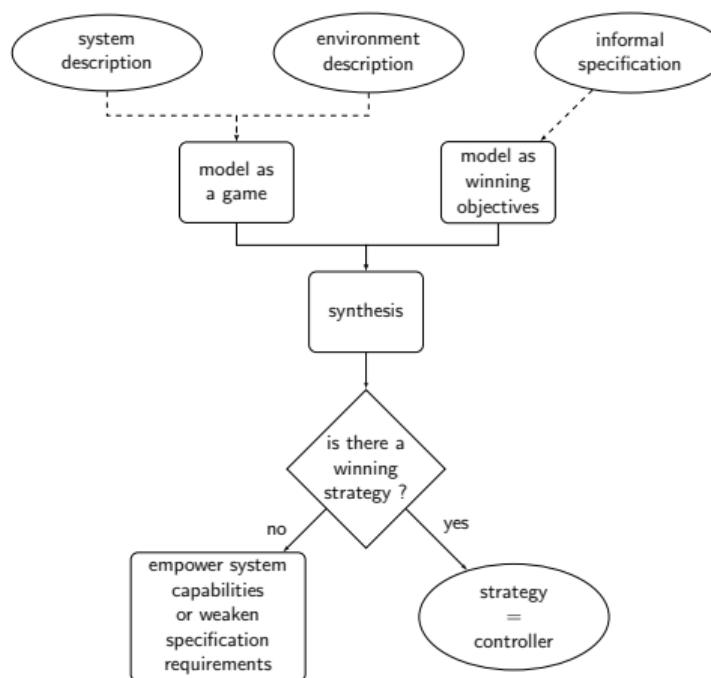
## 1 Controller synthesis

## 2 Memory

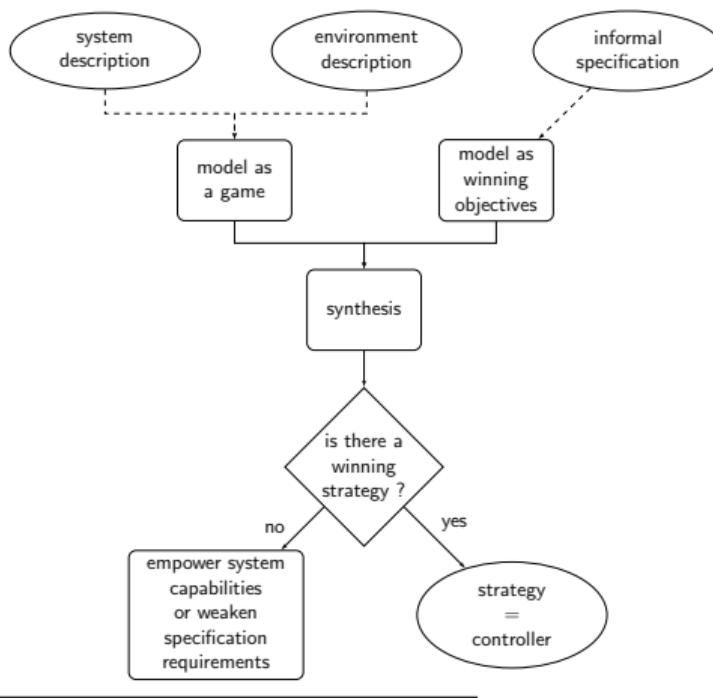
## 3 Randomness

## 4 Beyond Mealy machines

# Controller synthesis: a game-theoretic approach



# Controller synthesis: a game-theoretic approach

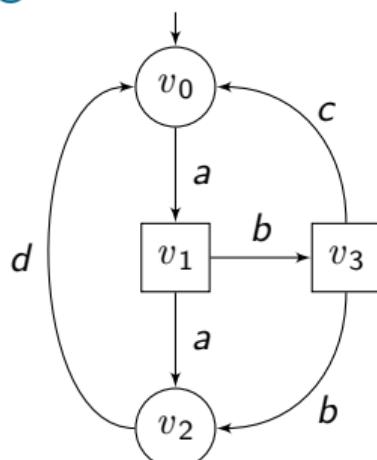


A plethora of models and objectives exist.<sup>1</sup>

Our focus here: how complex **strategies** need to be?

<sup>1</sup>Randour, "Automated Synthesis of Reliable and Efficient Systems Through Game Theory: A Case Study", 2013; Bloem, Chatterjee, and Jobstmann, "Graph Games and Reactive Synthesis", 2018; Fijalkow et al., Games on Graphs: From Logic and Automata to Algorithms, 2025.

## Two-player games



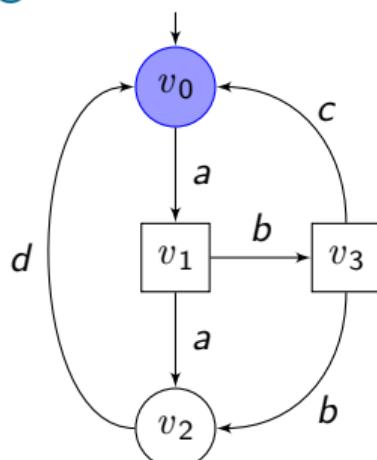
A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

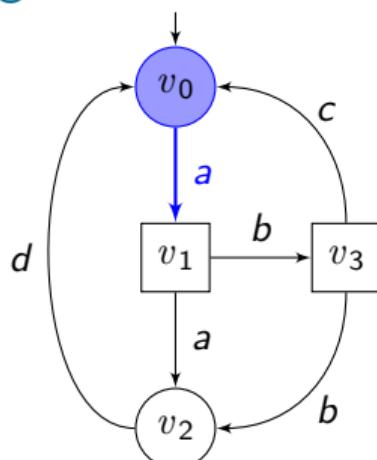
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play:

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

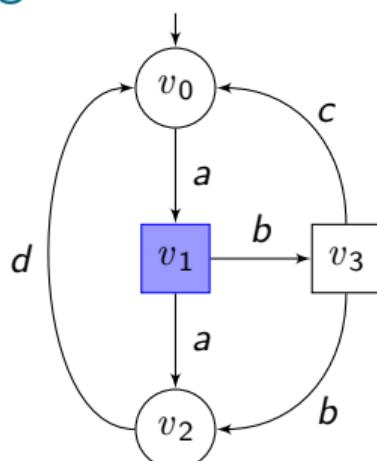
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play:  $a$

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

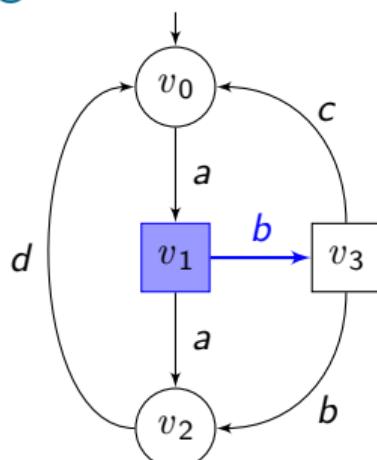
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play:  $a$

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

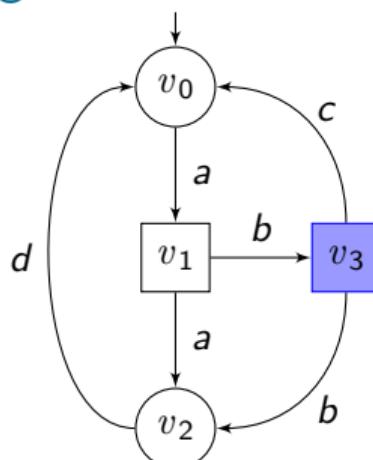
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *ab*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

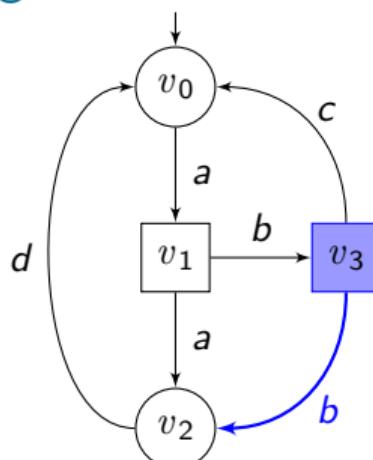
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *ab*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

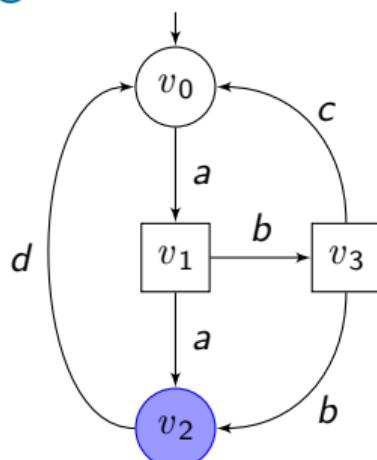
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *abb*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

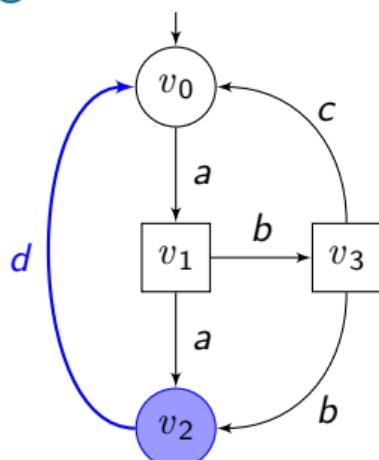
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *abb*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

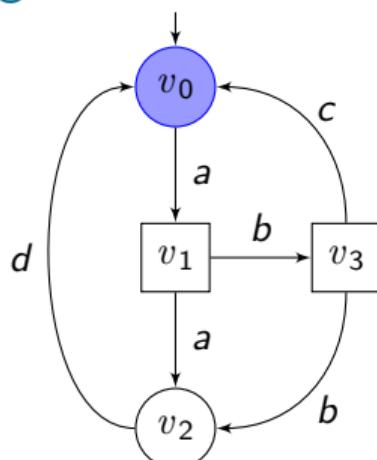
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *abbd*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

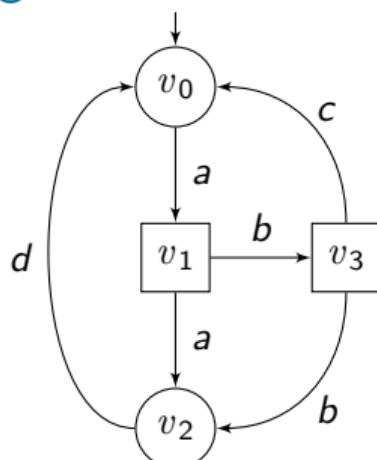
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play: *abbd*

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

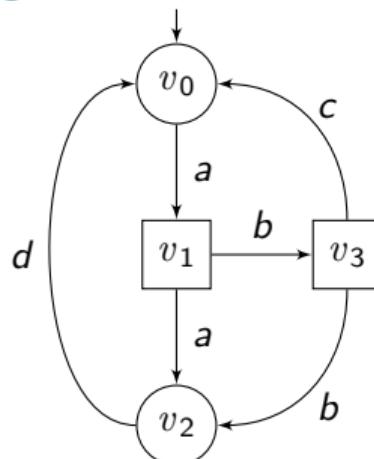
**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

Sample play:  $abbd\dots \in C^\omega$

## Two-player games



A two-player turn-based finite **arena**  
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$  with no deadlock.

**Color** function  $c: E \rightarrow C$ .

→ Players move a pebble along the edges creating an infinite **play**.

→ **Behavior of the system = sequence of colors.**

### Usual interpretation

$\mathcal{P}_\circlearrowleft$  (the system to control) tries to satisfy its **specification** while  $\mathcal{P}_\square$  (the environment) tries to prevent it from doing so.

# Specifications

They are encoded as some kind of *objective* defined using colors. Three main flavors:

# Specifications

They are encoded as some kind of *objective* defined using colors. Three main flavors:

- 1 A **winning condition**: a set of winning plays that  $\mathcal{P}_\circlearrowleft$  tries to realize. E.g.,  $\text{Reach}(t) = \{\pi = c_0c_1c_2 \dots \mid t \in \pi\}$ , for  $t \in C$  a given color, a *reachability* objective.

# Specifications

They are encoded as some kind of *objective* defined using colors. Three main flavors:

- 1 A **winning condition**: a set of winning plays that  $\mathcal{P}_\circlearrowleft$  tries to realize. E.g.,  $\text{Reach}(t) = \{\pi = c_0c_1c_2 \dots \mid t \in \pi\}$ , for  $t \in C$  a given color, a *reachability* objective.
- 2 A **payoff function** to optimize, assuming  $C \subset \mathbb{Q}$ . E.g., the *discounted sum* function, defined as  $\text{DS}(\pi) = \sum_{i=0}^{\infty} \gamma^i c_i$  for some discount factor  $\gamma \in (0, 1)$ .

# Specifications

They are encoded as some kind of *objective* defined using colors. Three main flavors:

- 1 A **winning condition**: a set of winning plays that  $\mathcal{P}_\circlearrowleft$  tries to realize. E.g.,  $\text{Reach}(t) = \{\pi = c_0c_1c_2 \dots \mid t \in \pi\}$ , for  $t \in C$  a given color, a *reachability* objective.
- 2 A **payoff function** to optimize, assuming  $C \subset \mathbb{Q}$ . E.g., the *discounted sum* function, defined as  $\text{DS}(\pi) = \sum_{i=0}^{\infty} \gamma^i c_i$  for some discount factor  $\gamma \in (0, 1)$ .
- 3 A **preference relation** defines a total preorder over sequences of colors, thus generalizing both previous concepts.

## Strategies

Player  $\mathcal{P}_\nabla$  chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (VE)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

## Strategies

Player  $\mathcal{P}_\nabla$  chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

↪ We are interested in the complexity of **optimal strategies**.

## Strategies

Player  $\mathcal{P}_\nabla$  chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

↪ We are interested in the complexity of **optimal strategies**.

Optimal strategies (using a preference relation  $\sqsubseteq$ )

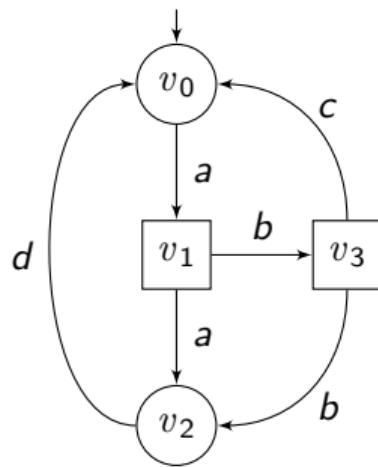
A strategy  $\sigma_\circlearrowleft$  of  $\mathcal{P}_\circlearrowleft$  is optimal if its **worst-case outcome** (i.e., considering all strategies of  $\mathcal{P}_\square$ ) is at least as good, with respect to  $\sqsubseteq$ , as that of any other strategy  $\sigma'_\circlearrowleft$ .

# MDPs & stochastic games

## Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.

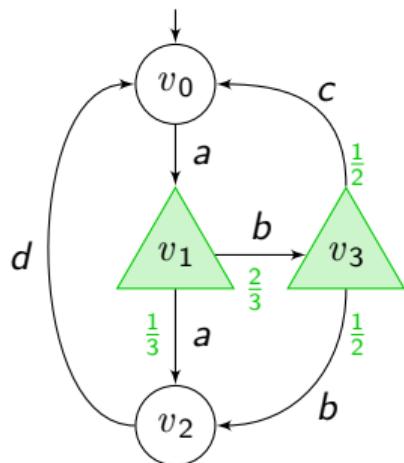
# MDPs & stochastic games



Two-player (deterministic) game.

$$V = V_{\circlearrowleft} \uplus V_{\square}.$$

# MDPs & stochastic games



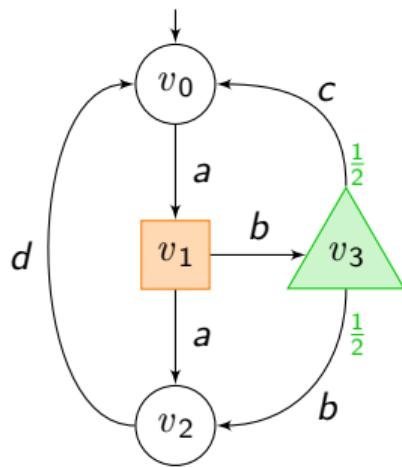
Markov decision process.

$$V = V_{\circlearrowleft} \uplus V_{\triangleleft}$$

Either  $\mathcal{P}_{\circlearrowleft}$  aims to maximize

- ▷  $\mathbb{P}^{\sigma_{\circlearrowleft}}[W]$  for some winning condition  $W$ ,
- ▷ or  $\mathbb{E}^{\sigma_{\circlearrowleft}}[f]$  for some payoff function  $f$ .

# MDPs & stochastic games



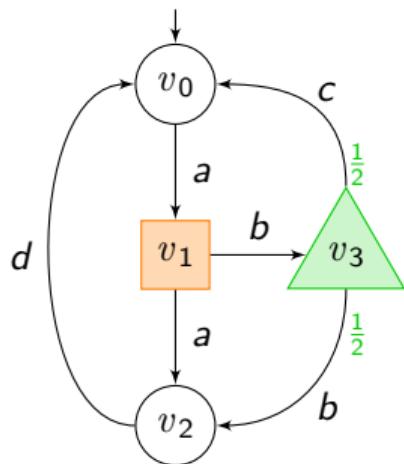
Stochastic game.

$$V = V_{\circlearrowleft} \uplus V_{\triangle} \uplus V_{\square}.$$

Either  $\mathcal{P}_{\circlearrowleft}$  aims to maximize, against the adversary  $\mathcal{P}_{\square}$ ,

- ▷  $\mathbb{P}^{\sigma_{\circlearrowleft}, \sigma_{\square}}[W]$  for some winning condition  $W$ ,
- ▷ or  $\mathbb{E}^{\sigma_{\circlearrowleft}, \sigma_{\square}}[f]$  for some payoff function  $f$ .

# MDPs & stochastic games



Stochastic game.

$$V = V_{\circ} \uplus V_{\Delta} \uplus V_{\square}.$$

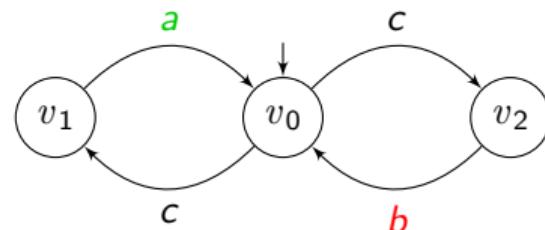
## Actions

We often use **actions** instead of stochastic vertices.

# Multiple objectives

## Combining objectives

Complex objectives arise when **combining** simple objectives, and usually require **more complex** strategies to play optimally.

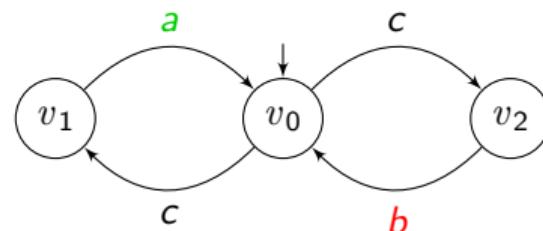


Seeing **a** and **b** infinitely often requires memory, but seeing only one does not (Büchi objective).

# Multiple objectives

## Combining objectives

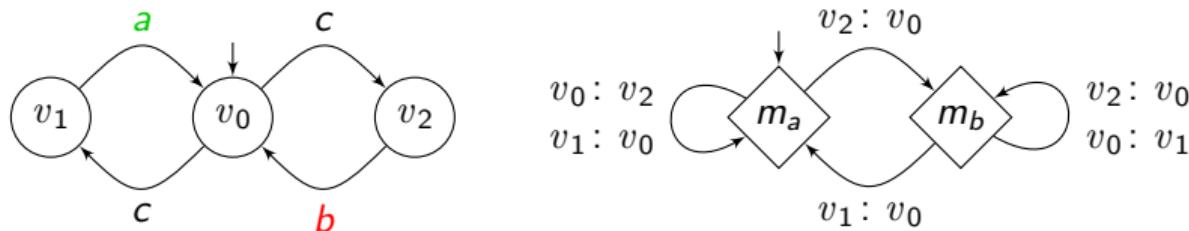
Complex objectives arise when **combining** simple objectives, and usually require **more complex** strategies to play optimally.



Seeing **a** and **b** infinitely often requires memory, but seeing only one does not (Büchi objective).

→ We are often interested in the **Pareto frontier**, i.e., all payoff vectors not dominated by another.

## Classical representation of strategies: Mealy machines

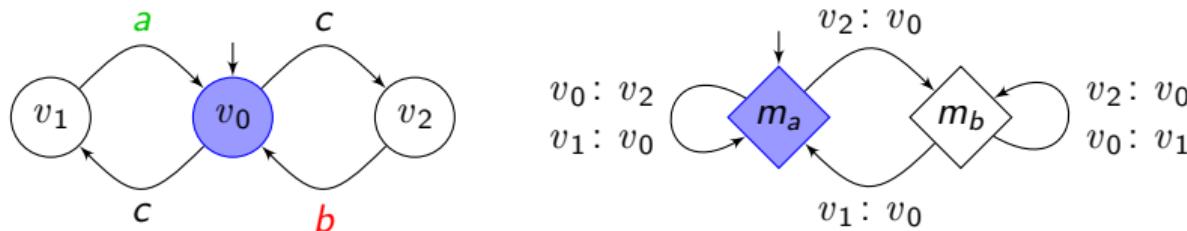


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

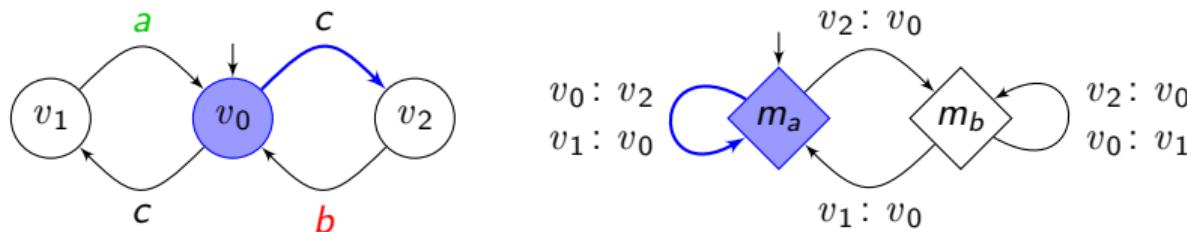


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

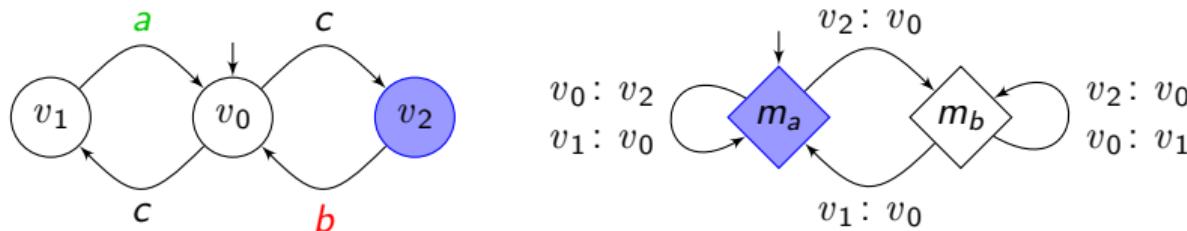


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

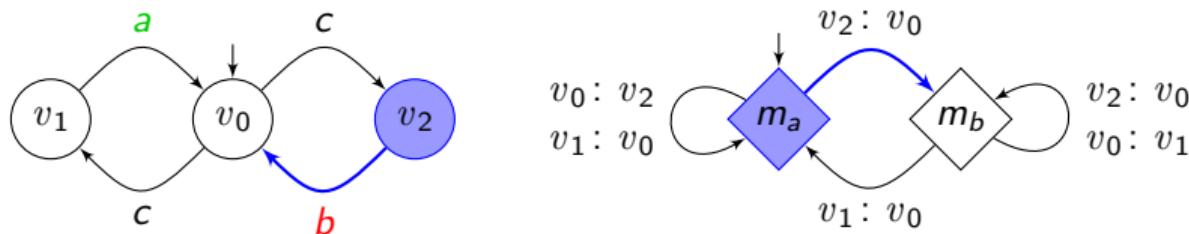


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

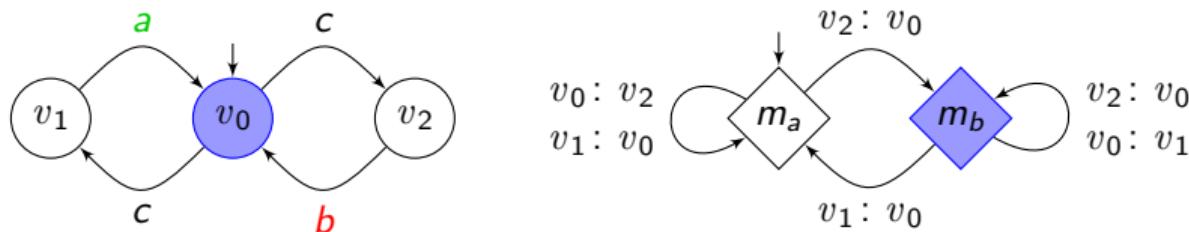


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

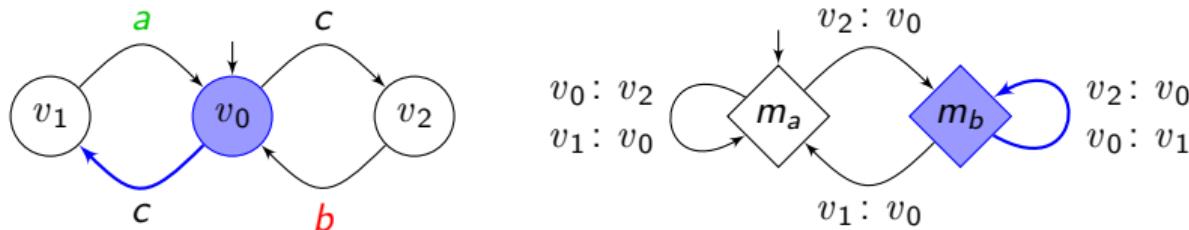


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

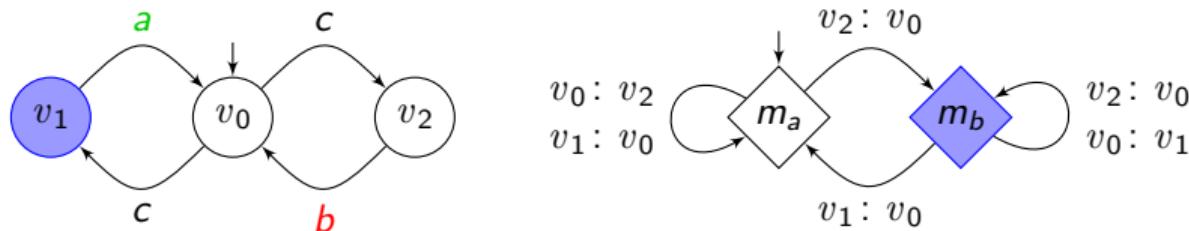


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

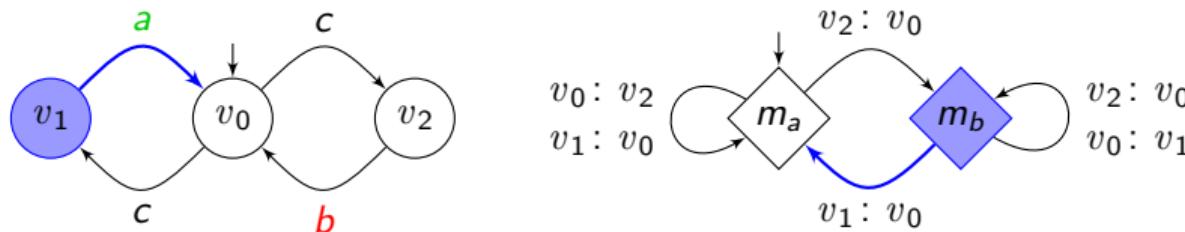


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

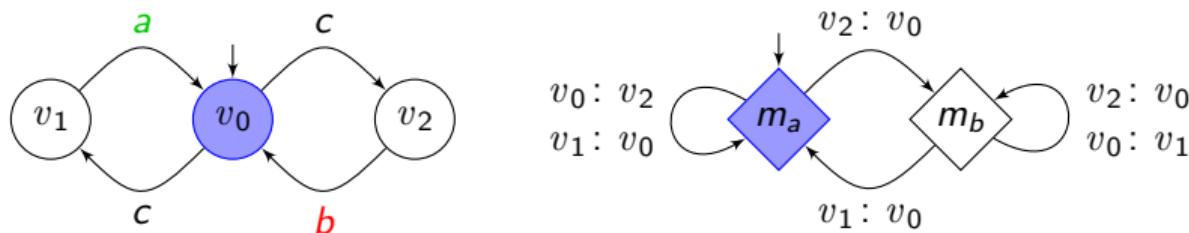


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

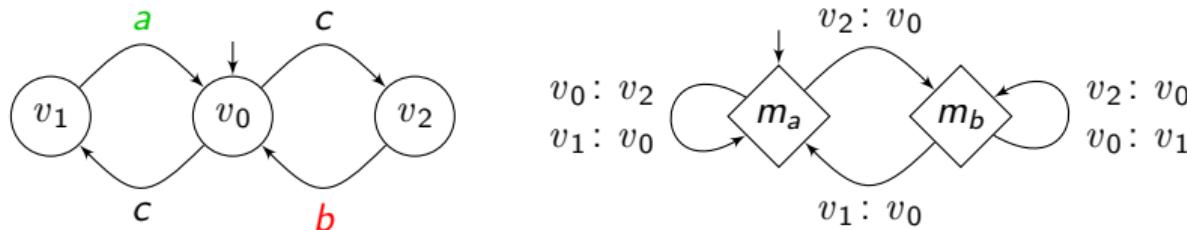


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## Classical representation of strategies: Mealy machines

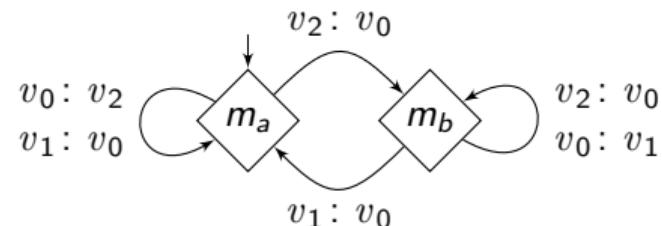


**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of *memory states*,
- ▷  $m_{\text{init}}$  is the *initial state*,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the *next-action function*,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the *update function*.

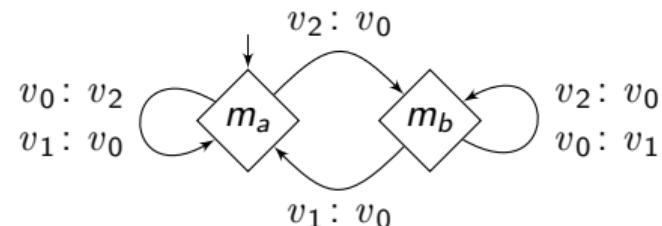
Finite memory if  $|M| < \infty$ , memoryless if  $|M| = 1$ .

## The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.

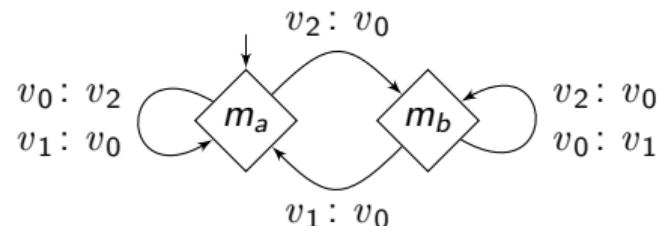
## The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.

Many other flavors exist: **chromatic** memory, with or without  **$\varepsilon$ -transitions**, with different types of **randomness**, etc.

## The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.

Many other flavors exist: **chromatic** memory, with or without  **$\varepsilon$ -transitions**, with different types of **randomness**, etc.

↪ We will discuss some of these.

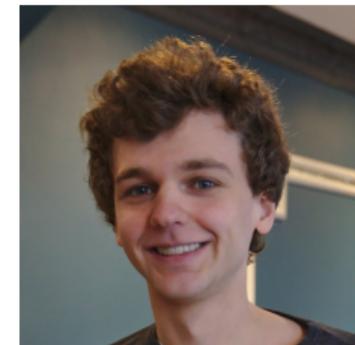
## 1 Controller synthesis

## 2 Memory

## 3 Randomness

## 4 Beyond Mealy machines

## Some amazing co-authors



Section mostly based on joint work with Patricia Bouyer, Stéphane Le Roux, Youssouf Oualhadj, and Pierre Vandenhove.<sup>2</sup>

---

<sup>2</sup>Bouyer, Le Roux, et al., "Games Where You Can Play Optimally with Arena-Independent Finite Memory", 2022; Bouyer, Oualhadj, et al., "Arena-Independent Finite-Memory Determinacy in Stochastic Games", 2023; Bouyer, Randour, and Vandenhove, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

## Memoryless strategies

Functions  $\sigma_\nabla: V_\nabla \rightarrow E$ .

- ▷ Equivalently, Mealy machines with one state.
- ▷ **Arguably**, the simplest kind of strategies.

## Memoryless strategies

Functions  $\sigma_\nabla : V_\nabla \rightarrow E$ .

- ▷ Equivalently, Mealy machines with one state.
- ▷ **Arguably**, the simplest kind of strategies.
- ▷ Sufficient to play optimally for most *single* objectives in (stochastic) games: reachability, parity, mean-payoff, discounted sum, etc.

## Starting point of our journey: *deterministic* games

### Gimbert and Zielonka's characterization<sup>3</sup>

Memoryless strategies suffice (for both players) for a preference relation  $\sqsubseteq$  iff  $\sqsubseteq$  and  $\sqsubseteq^{-1}$  are **monotone** and **selective**.

<sup>3</sup>Gimbert and Zielonka, "Games Where You Can Play Optimally Without Any Memory", 2005.

## Starting point of our journey: *deterministic* games

### Gimbert and Zielonka's characterization<sup>3</sup>

Memoryless strategies suffice (for both players) for a preference relation  $\sqsubseteq$  iff  $\sqsubseteq$  and  $\sqsubseteq^{-1}$  are **monotone** and **selective**.

### Corollary: one-to-two-player lift

If  $\sqsubseteq$  is such that

- 1 in all  $\mathcal{P}_\circlearrowleft$ -arenas,  $\mathcal{P}_\circlearrowleft$  has optimal memoryless strategies,
- 2 in all  $\mathcal{P}_\square$ -arenas,  $\mathcal{P}_\square$  has optimal memoryless strategies,

then **both** players have optimal memoryless strategies in all **two-player** arenas.

⇒ **Extremely useful as analyzing one-player games (i.e., graphs) is much easier.**

<sup>3</sup>Gimbert and Zielonka, "Games Where You Can Play Optimally Without Any Memory", 2005.

## Handling finite-memory strategies (1/3)

Why?

- More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).

## Handling finite-memory strategies (1/3)

Why?

- ▷ More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).
- ↪ One would hope for an equivalent of Gimbert and Zielonka's result for finite memory.

## Handling finite-memory strategies (1/3)

Why?

- ▷ More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).
- ↪ One would hope for an equivalent of Gimbert and Zielonka's result for finite memory.

Unfortunately, it does not hold.

## Handling finite-memory strategies (2/3)

Let  $C \subseteq \mathbb{Z}$  and the winning condition for  $\mathcal{P}_\circlearrowleft$  be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$

## Handling finite-memory strategies (2/3)

Let  $C \subseteq \mathbb{Z}$  and the winning condition for  $\mathcal{P}_\circlearrowleft$  be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$

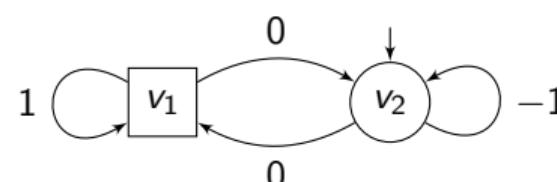
**Both one-player variants are finite-memory determined.**

## Handling finite-memory strategies (2/3)

Let  $C \subseteq \mathbb{Z}$  and the winning condition for  $\mathcal{P}_\circlearrowleft$  be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$

**Both one-player variants are finite-memory determined.**



**But the two-player one is not!**  
 $\Rightarrow \mathcal{P}_\circlearrowleft$  needs infinite memory to win.

## Handling finite-memory strategies (3/3)

A new frontier

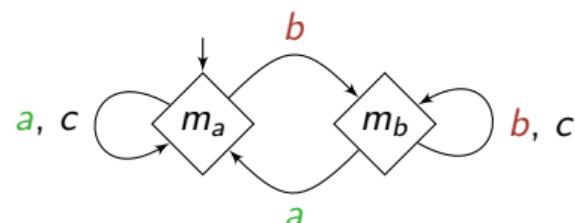
We focus on **arena-independent chromatic** memory structures.

## Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Example for  $C = \{a, b, c\}$  and objective Büchi( $a$ )  $\cap$  Büchi( $b$ ).

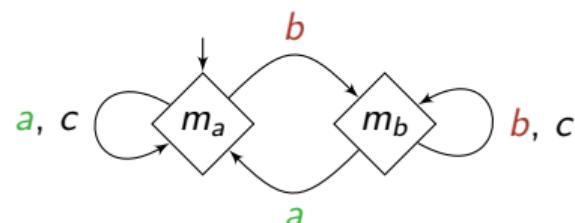


## Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Example for  $C = \{a, b, c\}$  and objective Büchi( $a$ )  $\cap$  Büchi( $b$ ).



This memory structure **suffices in all arenas**, i.e., it is always possible to find a suitable  $\alpha_{\text{nxt}}$  to build an optimal Mealy machine.

## Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Our characterization<sup>4</sup>

We obtain an equivalent to Gimbert and Zielonka's for finite memory:

- 1 a characterization through the concepts of  $\mathcal{M}$ -monotony and  $\mathcal{M}$ -selectivity,
- 2 a **one-to-two-player lift**.

<sup>4</sup>Bouyer, Le Roux, et al., "Games Where You Can Play Optimally with Arena-Independent Finite Memory", 2022.

## Extension to stochastic games

We lift<sup>5</sup> this result to **pure arena-independent finite-memory** strategies in **stochastic games**:

- 1 characterization based on generalizations of  $\mathcal{M}$ -monotony and  $\mathcal{M}$ -selectivity,
- 2 **one-to-two-player lift**, from MDPs to stochastic games.

---

<sup>5</sup>Bouyer, Oualhadj, et al., "Arena-Independent Finite-Memory Determinacy in Stochastic Games", 2023.

## Extension to infinite (deterministic) arenas (1/2)

We consider arenas of arbitrary cardinality and allow infinite branching.

### Observation

Memory requirements can be **higher in infinite arenas**: e.g., mean-payoff objectives require infinite memory.

## Extension to infinite (deterministic) arenas (1/2)

We consider arenas of arbitrary cardinality and allow infinite branching.

### Observation

Memory requirements can be **higher in infinite arenas**: e.g., mean-payoff objectives require infinite memory.

### The case of $\omega$ -regular objectives<sup>6</sup>

If a winning condition  $W$  is  $\omega$ -regular, **then** it admits finite-memory optimal strategies in all (infinite) arenas.

---

<sup>6</sup>Mostowski, "Regular expressions for infinite trees and a standard form of automata", 1985; Zielonka, "Infinite games on finitely coloured graphs with applications to automata on infinite trees", 1998.

## Extension to infinite (deterministic) arenas (2/2)

The converse<sup>7</sup>

If a **chromatic finite-memory** structure  $\mathcal{M}$  suffices for  $W$  in all infinite arenas, **then**  $W$  is  $\omega$ -regular.

→ We build a parity automaton for  $W$ , based on  $\mathcal{M}$  and  $\mathcal{S}_W$ , the *prefix-classifier* of  $W$  (recognizing its Myhill-Nerode classes).

<sup>7</sup>Bouyer, Randour, and Vandenbroucke, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

## Extension to infinite (deterministic) arenas (2/2)

The converse<sup>7</sup>

If a **chromatic finite-memory** structure  $\mathcal{M}$  suffices for  $W$  in all infinite arenas, **then**  $W$  is  $\omega$ -regular.

→ We build a parity automaton for  $W$ , based on  $\mathcal{M}$  and  $\mathcal{S}_W$ , the *prefix-classifier* of  $W$  (recognizing its Myhill-Nerode classes).

### Corollaries

- 1 Game-theoretical characterization of  $\omega$ -regularity.
- 2 **One-to-two-player lift** for infinite arenas.

<sup>7</sup>Bouyer, Randour, and Vandenbroucke, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

## Other criteria and characterizations

There is a plethora of results related to memory (models vary). Non-exhaustive list:

- ▷ characterizations through universal graphs,<sup>8</sup>
- ▷ tight memory bounds for sub-classes of objectives,<sup>9</sup>
- ▷ criteria for half-positionality,<sup>10</sup>
- ▷ one-to-multi-objective lift,<sup>11</sup>
- ▷ two-to-multi-player lift.<sup>12</sup>

→ Find more about **chromatic memory** in our survey.<sup>13</sup>

---

<sup>8</sup>Casares and Ohlmann, "Characterising memory in infinite games", 2025.

<sup>9</sup>Bouyer, Casares, et al., "Half-Positional Objectives Recognized by Deterministic Büchi Automata", 2024; Bouyer, Fijalkow, et al., "How to Play Optimally for Regular Objectives?", 2023; Casares and Ohlmann, "Positional  $\omega$ -regular languages", 2024.

<sup>10</sup>Kopczyński, "Half-positional Determinacy of Infinite Games", 2008.

<sup>11</sup>Le Roux, Pauly, and Randour, "Extending Finite-Memory Determinacy by Boolean Combination of Winning Conditions", 2018.

<sup>12</sup>Le Roux and Pauly, "Extending Finite Memory Determinacy to Multiplayer Games", 2016.

<sup>13</sup>Bouyer, Randour, and Vandenbroucke, "The True Colors of Memory: A Tour of Chromatic-Memory Strategies in Zero-Sum Games on Graphs", 2022.

## 1 Controller synthesis

## 2 Memory

## 3 Randomness

## 4 Beyond Mealy machines

# The amazing Mr. Main



Section mostly based on joint work with  
James C. A. Main.<sup>14</sup>

---

<sup>14</sup>Main and Randour, "Different Strokes in Randomised Strategies: Revisiting Kuhn's Theorem Under Finite-Memory Assumptions", 2024; Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

## Introducing randomness in strategies (1/2)

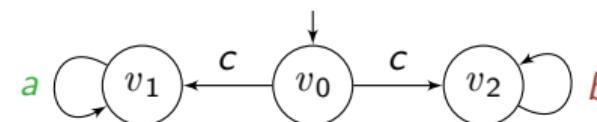
A **pure** strategy is a function  $\sigma_{\nabla}: (V E)^* V_{\nabla} \rightarrow E$ .

## Introducing randomness in strategies (1/2)

A **pure** strategy is a function  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ .

We may need **randomness** to deal with, e.g.,

- ▷ multiple objectives,
- ▷ concurrent games,
- ▷ imperfect information.



$$\text{Objective: } \mathbb{P}^{\sigma \circ} [\text{Reach}(a)] \geq \frac{1}{2} \wedge \mathbb{P}^{\sigma \circ} [\text{Reach}(b)] \geq \frac{1}{2}$$

↪ Achievable by tossing a coin in  $v_0$ .

## Introducing randomness in strategies (2/2)

Several ways of **randomizing**  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ :

## Introducing randomness in strategies (2/2)

Several ways of **randomizing**  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ :

### Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

## Introducing randomness in strategies (2/2)

Several ways of **randomizing**  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ :

**Behavioral strategies**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

**Mixed strategies**

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

## Introducing randomness in strategies (2/2)

Several ways of **randomizing**  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ :

### Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

### Mixed strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

### General strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E))$$

## Introducing randomness in strategies (2/2)

Several ways of **randomizing**  $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$ :

**Behavioral strategies**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

**Mixed strategies**

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

**General strategies**

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E))$$

Kuhn's theorem<sup>15</sup>

All three classes are equivalent in games of *perfect recall*.

→ Requires access to infinite memory and infinite support for distributions.

<sup>15</sup>Aumann, "Mixed and Behavior Strategies in Infinite Extensive Games", 1964; Bertrand, Genest, and Gimbert, "Qualitative Determinacy and Decidability of Stochastic Games with Signals", 2017.

## What about finite-memory strategies?

**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of memory states,
- ▷  $m_{\text{init}}$  is the initial state,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the next-action function,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the update function.

## What about finite-memory strategies?

**Mealy machine**  $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

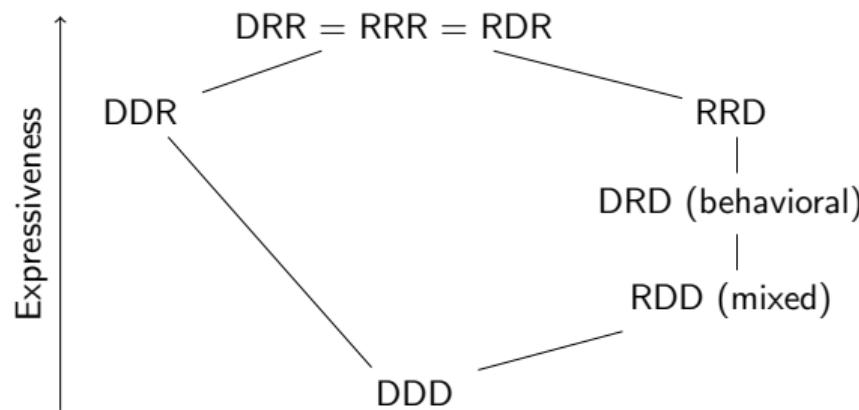
- ▷  $M$  is the set of memory states,
- ▷  $m_{\text{init}}$  is the initial state,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow E$  is the next-action function,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow M$  is the update function.

**Stochastic Mealy machine**  $\mathcal{M} = (M, \mu_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$ :

- ▷  $M$  is the set of memory states,
- ▷  $\mu_{\text{init}} \in \mathcal{D}(M)$  is the initial distribution,
- ▷  $\alpha_{\text{nxt}}: M \times V \rightarrow \mathcal{D}(E)$  is the next-action function,
- ▷  $\alpha_{\text{up}}: M \times E \rightarrow \mathcal{D}(M)$  is the update function.

⇒ **Three ways to add randomness: initialization, outputs, and updates.**

## Taxonomy<sup>16</sup> (1/2)



Classes **XYZ** with  $X, Y, Z \in \{D, R\}$ , where D stands for deterministic and R for random, and

- X characterizes the initialization,
- Y characterizes the next-action function,
- Z characterizes the update function.

<sup>16</sup>Main and Randour, "Different Strokes in Randomised Strategies: Revisiting Kuhn's Theorem Under Finite-Memory Assumptions", 2024.

## Taxonomy (2/2)

This taxonomy holds from one-player deterministic games (no collapse) up to concurrent partial-information multi-player games (equivalences hold).

## Taxonomy (2/2)

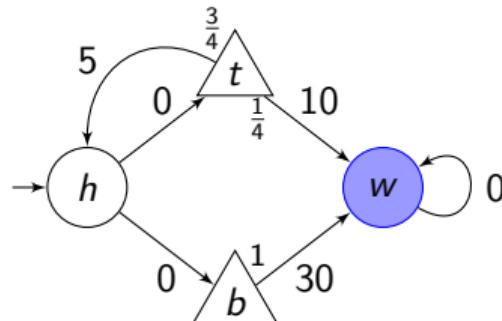
This taxonomy holds from one-player deterministic games (no collapse) up to concurrent partial-information multi-player games (equivalences hold).

↪ Collapses may arise for restricted classes of objectives (WiP).

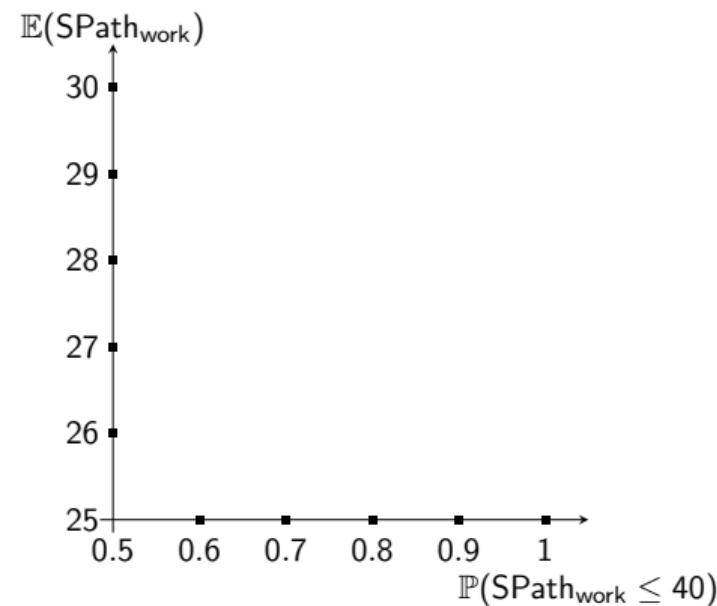
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



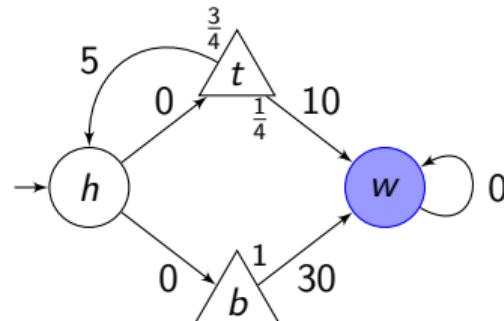
From *home*, take the *train* or *bike* to reach *work*.



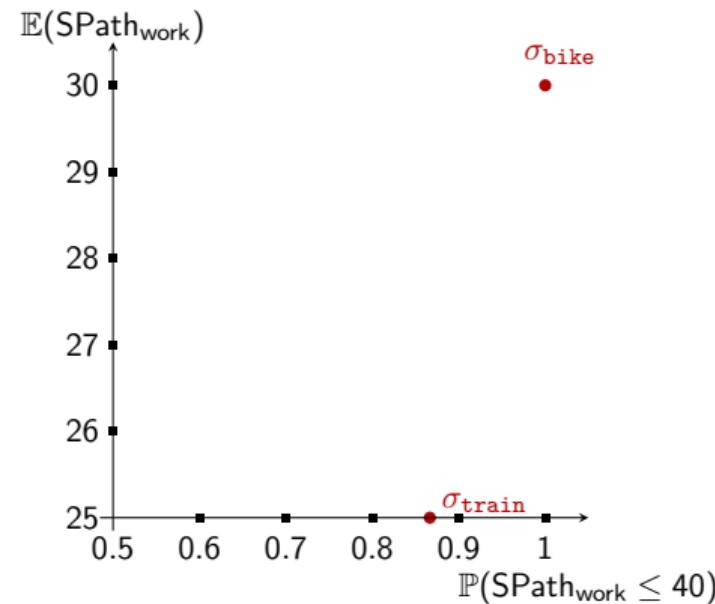
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



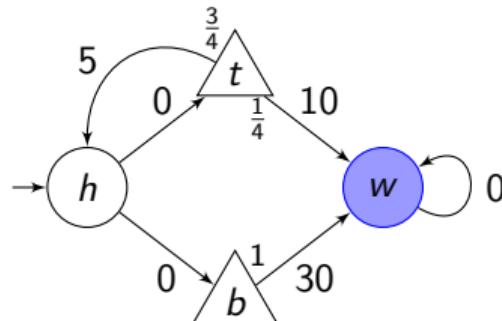
From *home*, take the *train* or *bike* to reach *work*.



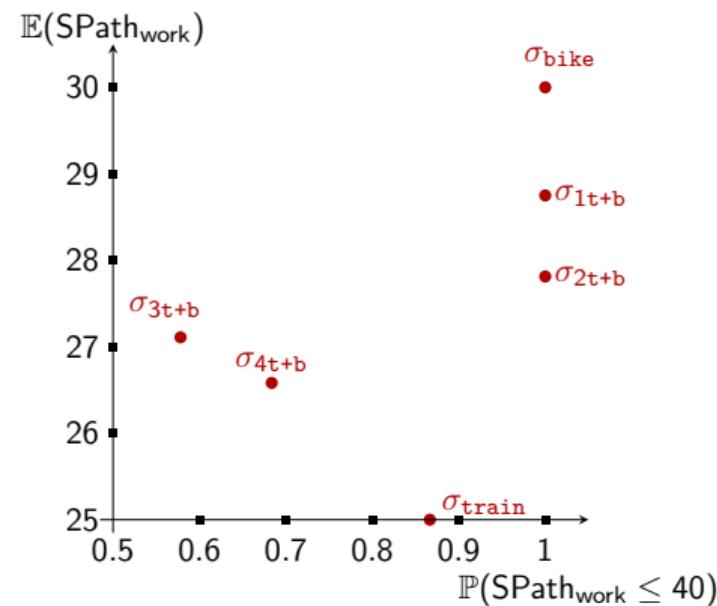
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



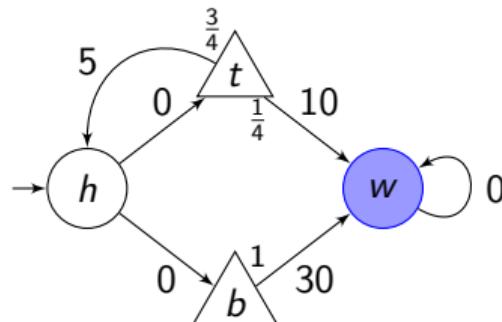
From *home*, take the *train* or *bike* to reach *work*.



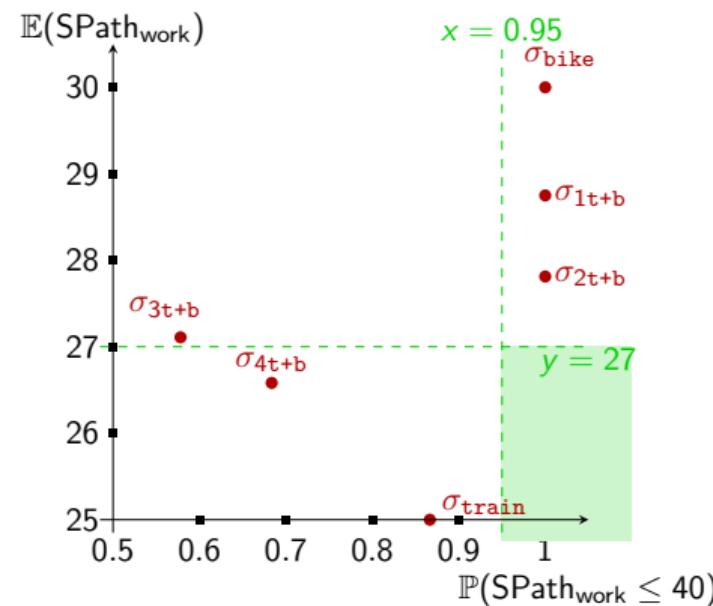
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



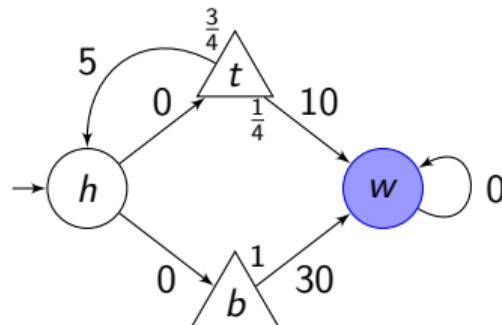
From *home*, take the *train* or *bike* to reach *work*.



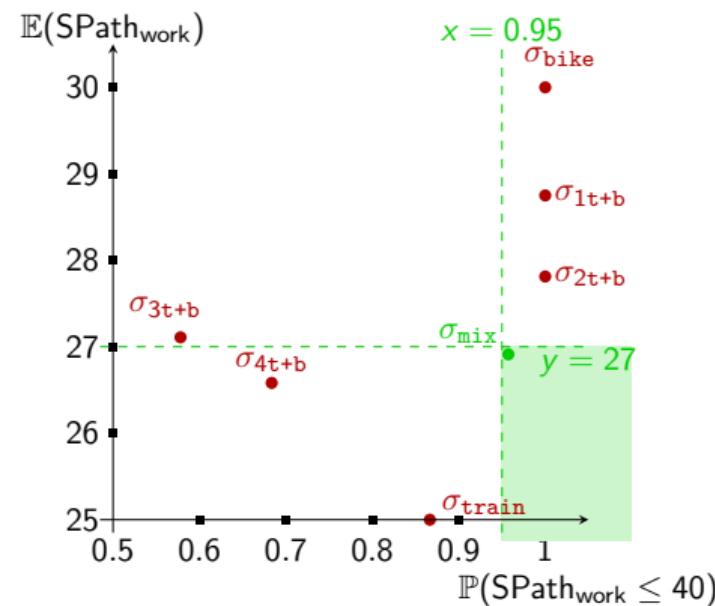
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



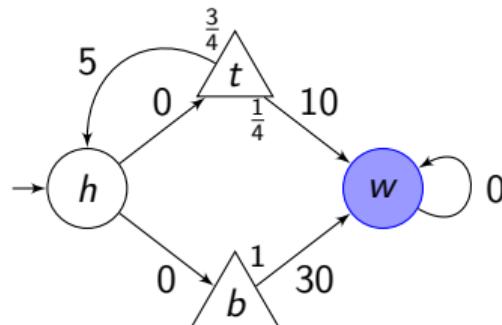
From *home*, take the *train* or *bike* to reach *work*.



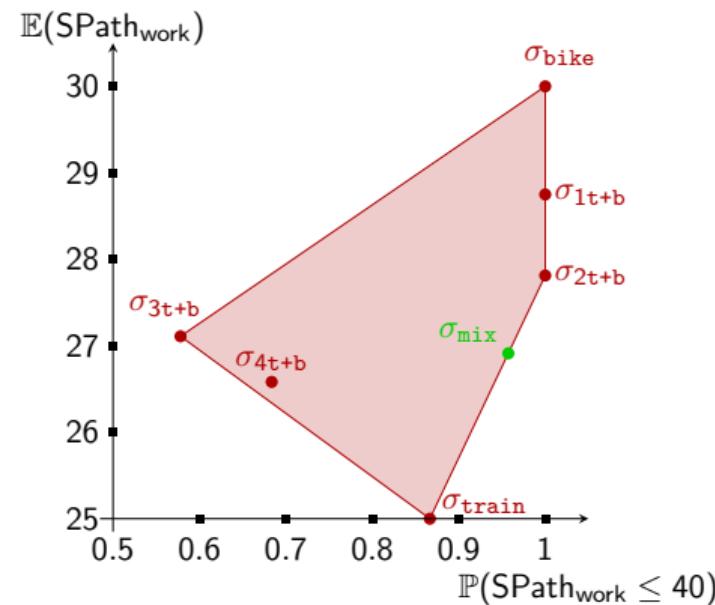
## Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



From *home*, take the *train* or *bike* to reach *work*.



## Multi-objectives MDPs (2/2)

We are interested in the structure of this **payoff set**.

Our result<sup>17</sup>

For *good* payoff functions ( $\sim$  expectancy is well-defined),

- 1 the set of achievable payoffs coincide with the convex hull of *pure* payoffs;
- 2 we can approximate *any* strategy  $\varepsilon$ -closely by **mixing** a bounded number of *pure* strategies.

<sup>17</sup> Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

## Multi-objectives MDPs (2/2)

We are interested in the structure of this **payoff set**.

Our result<sup>17</sup>

For *good* payoff functions ( $\sim$  expectancy is well-defined),

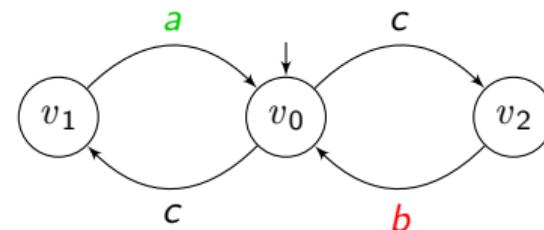
- 1 the set of achievable payoffs coincide with the convex hull of *pure* payoffs;
- 2 we can approximate *any* strategy  $\varepsilon$ -closely by **mixing** a bounded number of *pure* strategies.

⇒ **RDD-randomization is sufficient in most multi-objective MDPs.**

<sup>17</sup> Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

## Trading memory for randomness

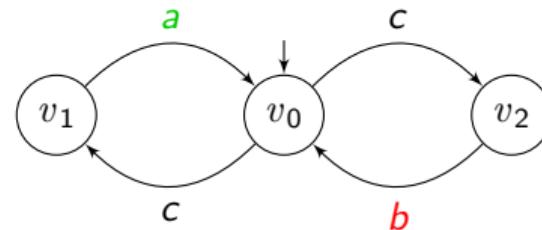
Recall this generalized Büchi game asking to see *a* and *b* infinitely often:



We need (a two-state) memory to win it with *pure* strategies.

## Trading memory for randomness

Recall this generalized Büchi game asking to see *a* and *b* infinitely often:

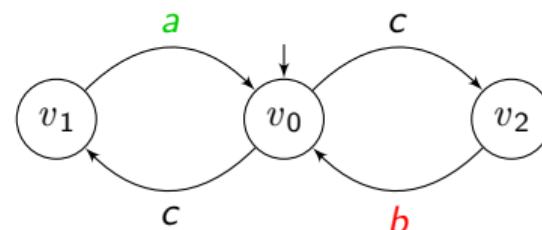


We need (a two-state) memory to win it with *pure* strategies.

**But a (behavioral) randomized memoryless strategy suffices to win with probability one:** playing  $v_1$  and  $v_2$  with non-zero probability ensures it.

## Trading memory for randomness

Recall this generalized Büchi game asking to see *a* and *b* infinitely often:



We need (a two-state) memory to win it with *pure* strategies.

But a (behavioral) randomized memoryless strategy suffices to win with probability one: playing  $v_1$  and  $v_2$  with non-zero probability ensures it.

↪ Memory can be traded for randomness for some classes of games/objectives.<sup>18</sup>

<sup>18</sup>Chatterjee, de Alfaro, and Henzinger, "Trading Memory for Randomness", 2004; Chatterjee, Randour, and Raskin, "Strategy synthesis for multi-dimensional quantitative objectives", 2014.

## 1 Controller synthesis

## 2 Memory

## 3 Randomness

## 4 Beyond Mealy machines

# An incomplete story

## Leitmotiv

Simpler strategies are better (for controller synthesis).

# An incomplete story

## Leitmotiv

Simpler strategies are better (for controller synthesis).

**But what is simple?**

# An incomplete story

## Leitmotiv

Simpler strategies are better (for controller synthesis).

**But what is simple?**

**Usual answer: small memory, no randomness.**

# An incomplete story

## Leitmotiv

Simpler strategies are better (for controller synthesis).

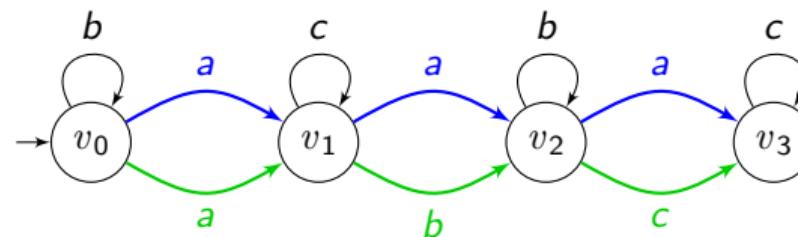
**But what is simple?**

**Usual answer: small memory, no randomness.**

↪ Let us question that.

## Not all memoryless strategies are created equal

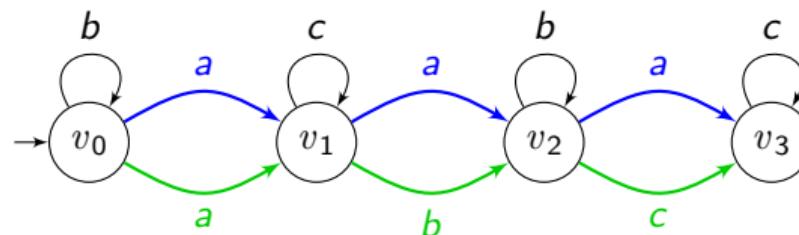
We want to reach  $v_3$ .



Intuitively, the blue strategy seems simpler than the green one.

## Not all memoryless strategies are created equal

We want to reach  $v_3$ .

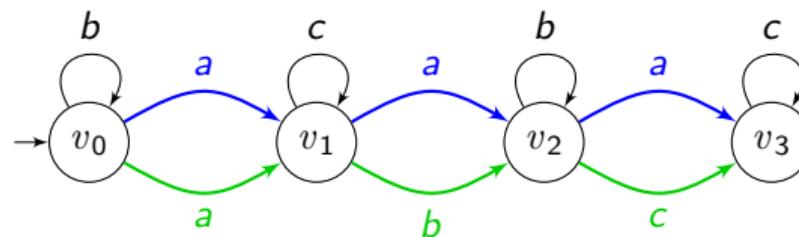


Intuitively, the **blue** strategy seems simpler than the **green** one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.

## Not all memoryless strategies are created equal

We want to reach  $v_3$ .

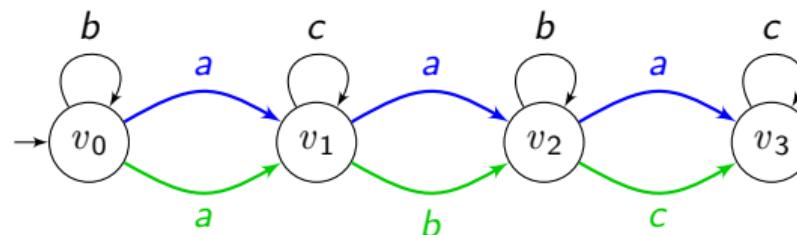


Intuitively, the **blue** strategy seems simpler than the **green** one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.
- ▷ The **representation of the next-action** function is mostly overlooked (basically a huge table).

## Not all memoryless strategies are created equal

We want to reach  $v_3$ .

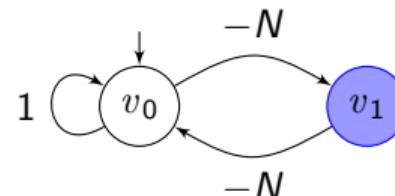


Intuitively, the **blue** strategy seems simpler than the **green** one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.
  - ▷ The **representation of the next-action** function is mostly overlooked (basically a huge table).
- **Memoryless strategies can already be too large to represent in practice!**

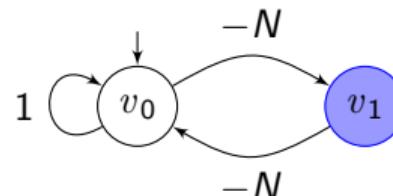
## Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., energy-Büchi objective with  $N \in \mathbb{N}$ .



## Memory is not always an issue

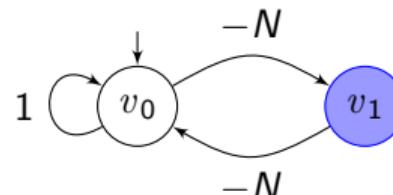
Multi-objectives games involving payoffs often require **exponential memory**. E.g., energy-Büchi objective with  $N \in \mathbb{N}$ .



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.

## Memory is not always an issue

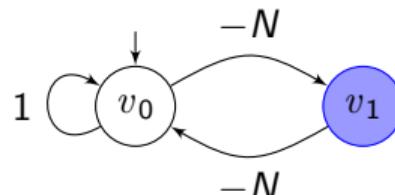
Multi-objectives games involving payoffs often require **exponential memory**. E.g., energy-Büchi objective with  $N \in \mathbb{N}$ .



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
  - **Polynomial representation if we allow the use of counters.**

## Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., energy-Büchi objective with  $N \in \mathbb{N}$ .



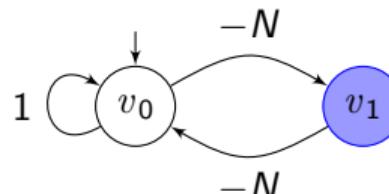
- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
  - **Polynomial representation if we allow the use of counters.**

### Hot take

We should explore novel notions of **simplicity**, and consider *alternative representations* of strategies/controllers.

## Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., energy-Büchi objective with  $N \in \mathbb{N}$ .



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
  - **Polynomial representation if we allow the use of counters.**

### Hot take

We should explore novel notions of **simplicity**, and consider *alternative representations* of strategies/controllers.

→ We quickly survey a few ones in the next slides.

# Structurally-enriched Mealy machines

Idea:

- ▷ Augment Mealy machines with **data structures**: e.g., counters.<sup>19</sup>
  - ▷ Avoid “flattening” structural information about the strategy: more succinct representations, better understandability, and closer to actual controllers.
- ⇒ **Changes our way of thinking which strategies are complex or not.**

---

<sup>19</sup> Blahoudek et al., “Qualitative Controller Synthesis for Consumption Markov Decision Processes”, 2020; Ajdarów et al., “Taming Infinity One Chunk at a Time: Concisely Represented Strategies in One-Counter MDPs”, 2025.

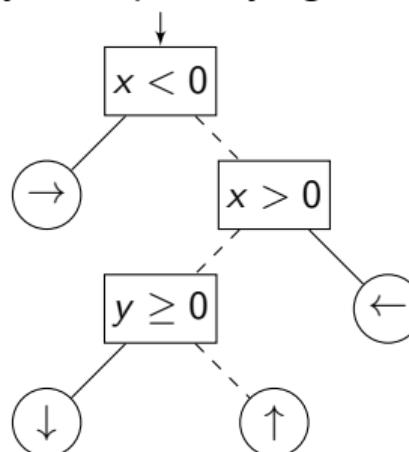
## Decision trees

- ▷ Structured state-space (e.g.,  $\subset \mathbb{Z}^n$ ) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

## Decision trees

- ▷ Structured state-space (e.g.,  $\subset \mathbb{Z}^n$ ) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

Toy example: trying to reach the center  $(0, 0)$  of a 2D-grid.



instead of

x	y	action
0	1	↓
0	2	↓
...	...	↓
-1	0	→
-1	1	→
...	...	...

## Decision trees

- ▷ Structured state-space (e.g.,  $\subset \mathbb{Z}^n$ ) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

Works well in practice...<sup>20</sup>

... starting from a given memoryless strategy.

---

<sup>20</sup>Brazdil, Chatterjee, Chmelik, et al., "Counterexample Explanation by Learning Small Strategies in Markov Decision Processes", 2015; Brazdil, Chatterjee, Kretinsky, et al., "Strategy Representation by Decision Trees in Reactive Synthesis", 2018.

# Other alternatives

## ■ Programmatic representations.

- ▷ Closer to realistic code, understandable.
- ▷ Strongly linked to the input format of the problem (e.g., PRISM code<sup>21</sup>), hard to generalize.

## ■ Models inspired by Turing machines.

- ▷ Powerful but hard to work with.
- ▷ Tentative notion of decision speed.<sup>22</sup>

## ■ Neural networks.

- ▷ Prevalent in RL.
- ▷ Hard to understand and verify.
- ▷ Can be coupled with finite-state-machine abstractions.<sup>23</sup>

---

<sup>23</sup>Shabadi, Fijalkow, and Matricon, "Programmatic Reinforcement Learning: Navigating Gridworlds", 2025.

<sup>23</sup>Gelderie, "Strategy machines: representation and complexity of strategies in infinite games", 2014.

<sup>23</sup>Carr, Jansen, and Topcu, "Verifiable RNN-Based Policies for POMDPs Under Temporal Logic Constraints", 2020.

# Wrap-up

## Focus

**Complexity of strategies** in controller synthesis.

# Wrap-up

## Focus

**Complexity of strategies** in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

# Wrap-up

## Focus

**Complexity of strategies** in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ **Many questions are still open!**

# Wrap-up

## Focus

**Complexity of strategies** in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ Many questions are still open!

**Strategy complexity  $\neq$  representation complexity.**

# Wrap-up

## Focus

**Complexity of strategies** in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ Many questions are still open!

Strategy complexity  $\neq$  representation complexity.

## Take-home message

We need a **proper theory of complexity**, and a **toolbox of different representations**.

→ Ongoing project **ContolleRS**.

# Thank you! Any question?