

Simplicity Lies in the Eye of the Beholder

A Strategic Perspective on Controllers in Reactive Synthesis

Mickael Randour

F.R.S.-FNRS & UMONS – Université de Mons, Belgium

October 3, 2025

Reachability Problems 2025



Special thanks to the *Delegación General Valonia-Bruselas en España*.

The talk in one slide

The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

Simpler is better:

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

Simpler is better:

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

Simpler is better:

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

But how to define complexity and how to measure it?

The talk in one slide

Strategies = **formal blueprints** for real-world controllers.

Simpler is better:

- ▷ easier to understand,
- ▷ cheaper to produce and maintain.

Aim of this survey talk

Understanding **how complex** strategies need to be.

But how to define complexity and how to measure it?

↪ **That is our topic of the today.**

The talk in one **more** slide

The talk in one **more** slide

Yes, I lied, and I will lie even more. The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

The talk in one **more** slide

Yes, I lied, and I will lie even more. The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

There will be some hand-waving and approximations to keep the talk high level.

The talk in one **more** slide

Yes, I lied, and I will lie even more. The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

There will be some hand-waving and approximations to keep the talk high level.

→ **I will focus on recent work with marvelous co-authors.**

The talk in one **more** slide

Yes, I lied, and I will lie even more. The results I will survey span numerous combinations of

- ▷ game models,
- ▷ strategy models,
- ▷ objectives,
- ▷ decision problems...

There will be some hand-waving and approximations to keep the talk high level.

→ **I will focus on recent work with marvelous co-authors.**

→ Check the paper for more details.

1 Controller synthesis

2 Memory

3 Randomness

4 Beyond Mealy machines


1 Controller synthesis

2 Memory


3 Randomness

4 Beyond Mealy machines

Controller synthesis: a game-theoretic approach



Controller synthesis: a game-theoretic approach




A plethora of models and objectives exist.¹

Our focus here: how complex **strategies** need to be?

¹ Randour, "Automated Synthesis of Reliable and Efficient Systems Through Game Theory: A Case Study", 2013; Bloem, Chatterjee, and Jobstmann, "Graph Games and Reactive Synthesis", 2018; Fijalkow et al., Games on Graphs: From Logic and Automata to Algorithms, 2025.

Two-player games




A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.

Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play:

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: a

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: a

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *ab*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *ab*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *abb*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *abb*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *abbd*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: *abbd*

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_\circlearrowleft, V_\square, E)$ with no deadlock.


Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Sample play: $abbd\dots \in C^\omega$

Two-player games



A two-player turn-based finite **arena**
 $\mathcal{A} = (V_{\circlearrowleft}, V_{\square}, E)$ with no deadlock.

Color function $c: E \rightarrow C$.

→ Players move a pebble along the edges creating an infinite **play**.

↪ Behavior of the system = sequence of colors.

Usual interpretation

$\mathcal{P}_{\circlearrowleft}$ (the system to control) tries to satisfy its **specification** while \mathcal{P}_{\square} (the environment) tries to prevent it from doing so.

Specifications

They are encoded as some kind of *objective* defined using colors.
Three main flavors:

Specifications

They are encoded as some kind of *objective* defined using colors.

Three main flavors:

- 1 A **winning condition**: a set of winning plays that $\mathcal{P}_\circlearrowleft$ tries to realize. E.g., $\text{Reach}(t) = \{\pi = c_0 c_1 c_2 \dots \mid t \in \pi\}$, for $t \in C$ a given color, a *reachability* objective.

Specifications

They are encoded as some kind of *objective* defined using colors.

Three main flavors:

- 1 A **winning condition**: a set of winning plays that $\mathcal{P}_\circlearrowleft$ tries to realize. E.g., $\text{Reach}(t) = \{\pi = c_0 c_1 c_2 \dots \mid t \in \pi\}$, for $t \in C$ a given color, a *reachability* objective.
- 2 A **payoff function** to optimize, assuming $C \subset \mathbb{Q}$. E.g., the *discounted sum* function, defined as $\text{DS}(\pi) = \sum_{i=0}^{\infty} \gamma^i c_i$ for some discount factor $\gamma \in (0, 1)$.

Specifications

They are encoded as some kind of *objective* defined using colors.

Three main flavors:

- 1 A **winning condition**: a set of winning plays that $\mathcal{P}_\circlearrowleft$ tries to realize. E.g., $\text{Reach}(t) = \{\pi = c_0 c_1 c_2 \dots \mid t \in \pi\}$, for $t \in C$ a given color, a *reachability* objective.
- 2 A **payoff function** to optimize, assuming $C \subset \mathbb{Q}$. E.g., the *discounted sum* function, defined as $\text{DS}(\pi) = \sum_{i=0}^{\infty} \gamma^i c_i$ for some discount factor $\gamma \in (0, 1)$.
- 3 A **preference relation** defines a total preorder over sequences of colors, thus generalizing both previous concepts.

Strategies

Player \mathcal{P}_∇ chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (VE)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

Strategies

Player \mathcal{P}_∇ chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

↪ We are interested in the complexity of **optimal strategies**.

Strategies

Player \mathcal{P}_∇ chooses outgoing edges following a **strategy**

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$$

consistent with the underlying graph.

→ We are interested in the complexity of **optimal strategies**.

Optimal strategies (using a preference relation \sqsubseteq)

A strategy σ_\circlearrowleft of $\mathcal{P}_\circlearrowleft$ is optimal if its **worst-case outcome** (i.e., considering all strategies of \mathcal{P}_\square) is at least as good, with respect to \sqsubseteq , as that of any other strategy σ'_\circlearrowleft .

MDPs & stochastic games


Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.

MDPs & stochastic games

Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.




Two-player (deterministic) game.

$$V = V_{\circlearrowleft} \uplus V_{\square}.$$

MDPs & stochastic games

Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.



Markov decision process.

$$V = V_{\circlearrowleft} \uplus V_{\triangleleft}.$$


Either $\mathcal{P}_{\circlearrowleft}$ aims to maximize

- ▷ $\mathbb{P}^{\sigma_{\circlearrowleft}}[W]$ for some winning condition W ,
- ▷ or $\mathbb{E}^{\sigma_{\circlearrowleft}}[f]$ for some payoff function f .

MDPs & stochastic games

Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.



Stochastic game.

$$V = V_{\circ} \uplus V_{\Delta} \uplus V_{\square}.$$


Either \mathcal{P}_{\circ} aims to maximize, against the adversary \mathcal{P}_{\square} ,

- ▷ $\mathbb{P}^{\sigma_{\circ}, \sigma_{\square}}[W]$ for some winning condition W ,
- ▷ or $\mathbb{E}^{\sigma_{\circ}, \sigma_{\square}}[f]$ for some payoff function f .

MDPs & stochastic games

Why?

In many real-world scenarios, the environment is not fully antagonistic, but exhibits stochastic behaviors.



Stochastic game.

$$V = V_{\circ} \uplus V_{\Delta} \uplus V_{\square}.$$


Actions

We often use **actions** instead of stochastic vertices.

Multiple objectives

Combining objectives

Complex objectives arise when **combining** simple objectives, and usually require **more complex** strategies to play optimally.




Seeing **a** and **b** infinitely often requires memory, but seeing only one does not (Büchi objective).

Multiple objectives

Combining objectives



Complex objectives arise when **combining** simple objectives, and usually require **more complex** strategies to play optimally.



Seeing **a** and **b** infinitely often requires memory, but seeing only one does not (Büchi objective).

→ We are often interested in the **Pareto frontier**, i.e., all payoff vectors not dominated by another.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines





Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.

Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

Classical representation of strategies: Mealy machines




Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of *memory states*,
- ▷ m_{init} is the *initial state*,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the *next-action function*,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the *update function*.


Finite memory if $|M| < \infty$, memoryless if $|M| = 1$.

The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.


The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.

Many other flavors exist: **chromatic** memory, with or without **ϵ -transitions**, with different types of **randomness**, etc.

The ice cream conundrum



This Mealy machine uses **chaotic** (or general) memory: it looks at the actual vertices of the game to update its memory.

Many other flavors exist: **chromatic** memory, with or without **ϵ -transitions**, with different types of **randomness**, etc.

→ We will discuss some of these.


1 Controller synthesis

2 Memory

3 Randomness

4 Beyond Mealy machines

Some amazing co-authors



Section mostly based on joint work with Patricia Bouyer, Stéphane Le Roux, Youssouf Oualhadj, and Pierre Vandenhove.²

²Bouyer, Le Roux, et al., "Games Where You Can Play Optimally with Arena-Independent Finite Memory", 2022; Bouyer, Oualhadj, et al., "Arena-Independent Finite-Memory Determinacy in Stochastic Games", 2023; Bouyer, Randour, and Vandenhove, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

Memoryless strategies

Functions $\sigma_\nabla: V_\nabla \rightarrow E$.

- ▷ Equivalently, Mealy machines with one state.
- ▷ **Arguably**, the simplest kind of strategies.

Memoryless strategies

Functions $\sigma_\nabla: V_\nabla \rightarrow E$.

- ▷ Equivalently, Mealy machines with one state.
- ▷ **Arguably**, the simplest kind of strategies.
- ▷ Sufficient to play optimally for most *single* objectives in (stochastic) games: reachability, parity, mean-payoff, discounted sum, etc.

Starting point of our journey: *deterministic* games

Gimbert and Zielonka's characterization³

Memoryless strategies suffice (for both players) for a preference relation \sqsubseteq iff \sqsubseteq and \sqsubseteq^{-1} are **monotone** and **selective**.

³Gimbert and Zielonka, "Games Where You Can Play Optimally Without Any Memory", 2005.

Starting point of our journey: *deterministic* games

Gimbert and Zielonka's characterization³

Memoryless strategies suffice (for both players) for a preference relation \sqsubseteq iff \sqsubseteq and \sqsubseteq^{-1} are **monotone** and **selective**.

Corollary: one-to-two-player lift

If \sqsubseteq is such that

- 1 in all \mathcal{P}_\circ -arenas, \mathcal{P}_\circ has optimal memoryless strategies,
 - 2 in all \mathcal{P}_\square -arenas, \mathcal{P}_\square has optimal memoryless strategies,
- then **both** players have optimal memoryless strategies in all **two-player** arenas.

⇒ **Extremely useful as analyzing one-player games (i.e., graphs) is much easier.**

³Gimbert and Zielonka, "Games Where You Can Play Optimally Without Any Memory", 2005.

Handling finite-memory strategies (1/3)

Why?

- More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).

Handling finite-memory strategies (1/3)

Why?

- ▷ More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).

- ↪ One would hope for an equivalent of Gimbert and Zielonka's result for finite memory.

Handling finite-memory strategies (1/3)

Why?

- ▷ More complex objectives may require **finite** (multi-Büchi) or **infinite** memory (multi-mean-payoff).

- ↪ One would hope for an equivalent of Gimbert and Zielonka's result for finite memory.

Unfortunately, it does not hold.

Handling finite-memory strategies (2/3)

Let $C \subseteq \mathbb{Z}$ and the winning condition for $\mathcal{P}_\circlearrowleft$ be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$

Handling finite-memory strategies (2/3)

Let $C \subseteq \mathbb{Z}$ and the winning condition for $\mathcal{P}_\circlearrowleft$ be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$


Both one-player variants are finite-memory determined.

Handling finite-memory strategies (2/3)

Let $C \subseteq \mathbb{Z}$ and the winning condition for $\mathcal{P}_\circlearrowleft$ be

$$\overline{TP}(\pi) = \infty \quad \vee \quad \exists^\infty n \in \mathbb{N}, \sum_{i=0}^n c_i = 0$$

Both one-player variants are finite-memory determined.



But the two-player one is not!
 $\Rightarrow \mathcal{P}_\circlearrowleft$ needs infinite memory to win.

Handling finite-memory strategies (3/3)

A new frontier


We focus on **arena-independent chromatic** memory structures.

Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Example for $C = \{a, b, c\}$ and objective Büchi(a) \cap Büchi(b).




Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Example for $C = \{a, b, c\}$ and objective $\text{Büchi}(a) \cap \text{Büchi}(b)$.



This memory structure **suffices in all arenas**, i.e., it is always possible to find a suitable α_{nxt} to build an optimal Mealy machine.

Handling finite-memory strategies (3/3)

A new frontier

We focus on **arena-independent chromatic** memory structures.

Our characterization⁴

We obtain an equivalent to Gimbert and Zielonka's for finite memory:

- 1 a characterization through the concepts of **M -monotony** and **M -selectivity**,
- 2 a **one-to-two-player lift**.

⁴ Bouyer, Le Roux, et al., "Games Where You Can Play Optimally with Arena-Independent Finite Memory", 2022.

Extension to stochastic games

We lift⁵ this result to **pure arena-independent finite-memory** strategies in **stochastic games**:

- 1 characterization based on generalizations of \mathcal{M} -monotony and \mathcal{M} -selectivity,
- 2 **one-to-two-player lift**, from MDPs to stochastic games.

⁵Bouyer, Oualhadj, et al., "Arena-Independent Finite-Memory Determinacy in Stochastic Games", 2023.

Extension to infinite (deterministic) arenas (1/2)

We consider arenas of arbitrary cardinality and allow infinite branching.

Observation

Memory requirements can be **higher in infinite arenas**: e.g., mean-payoff objectives require infinite memory.

Extension to infinite (deterministic) arenas (1/2)

We consider arenas of arbitrary cardinality and allow infinite branching.

Observation

Memory requirements can be **higher in infinite arenas**: e.g., mean-payoff objectives require infinite memory.

The case of ω -regular objectives⁶

If a winning condition W is ω -regular, **then** it admits finite-memory optimal strategies in all (infinite) arenas.

⁶Mostowski, "Regular expressions for infinite trees and a standard form of automata", 1985; Zielonka, "Infinite games on finitely coloured graphs with applications to automata on infinite trees", 1998.

Extension to infinite (deterministic) arenas (2/2)

The converse⁷

If a **chromatic finite-memory** structure \mathcal{M} suffices for W in all infinite arenas, **then** W is ω -regular.

→ We build a parity automaton for W , based on \mathcal{M} and \mathcal{S}_W , the **prefix-classifier** of W (recognizing its Myhill-Nerode classes).

⁷Bouyer, Randour, and Vandenbroucke, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

Extension to infinite (deterministic) arenas (2/2)

The converse⁷

If a **chromatic finite-memory** structure \mathcal{M} suffices for W in all infinite arenas, **then** W is ω -regular.

→ We build a parity automaton for W , based on \mathcal{M} and \mathcal{S}_W , the **prefix-classifier** of W (recognizing its Myhill-Nerode classes).

Corollaries

- 1 Game-theoretical characterization of ω -regularity.
- 2 **One-to-two-player lift** for infinite arenas.

⁷Bouyer, Randour, and Vandenbroucke, "Characterizing Omega-Regularity through Finite-Memory Determinacy of Games on Infinite Graphs", 2023.

Other criteria and characterizations

There is a plethora of results related to memory (models vary).

Non-exhaustive list:

- ▷ characterizations through universal graphs,⁸
- ▷ tight memory bounds for sub-classes of objectives,⁹
- ▷ criteria for half-positionality,¹⁰
- ▷ one-to-multi-objective lift,¹¹
- ▷ two-to-multi-player lift.¹²

↪ Find more about **chromatic memory** in our survey.¹³

⁸ Casares and Ohlmann, "Characterising memory in infinite games", 2025.

⁹ Bouyer, Casares, et al., "Half-Positional Objectives Recognized by Deterministic Büchi Automata", 2024; Bouyer, Fijalkow, et al., "How to Play Optimally for Regular Objectives?", 2023; Casares and Ohlmann, "Positional ω -regular languages", 2024.

¹⁰ Kopczyński, "Half-positional Determinacy of Infinite Games", 2008.

¹¹ Le Roux, Pauly, and Randour, "Extending Finite-Memory Determinacy by Boolean Combination of Winning Conditions", 2018.

¹² Le Roux and Pauly, "Extending Finite Memory Determinacy to Multiplayer Games", 2016.

¹³ Bouyer, Randour, and Vandenhove, "The True Colors of Memory: A Tour of Chromatic-Memory Strategies in Zero-Sum Games on Graphs", 2022.

1 Controller synthesis

2 Memory

3 Randomness

4 Beyond Mealy machines

The amazing Mr. Main



Section mostly based on joint work with
James C. A. Main.¹⁴

¹⁴ Main and Randour, "Different Strokes in Randomised Strategies: Revisiting Kuhn's Theorem Under Finite-Memory Assumptions", 2024; Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

Introducing randomness in strategies (1/2)


A **pure** strategy is a function $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$.

Introducing randomness in strategies (1/2)

A **pure** strategy is a function $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$.

We may need **randomness** to deal with, e.g.,

- ▷ multiple objectives,
- ▷ concurrent games,
- ▷ imperfect information.



Objective: $\mathbb{P}^{\sigma_0}[\text{Reach}(a)] \geq \frac{1}{2} \wedge \mathbb{P}^{\sigma_0}[\text{Reach}(b)] \geq \frac{1}{2}$

→ **Achievable by tossing a coin in v_0 .**

Introducing randomness in strategies (2/2)

Several ways of **randomizing** $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$:

Introducing randomness in strategies (2/2)

Several ways of **randomizing** $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$:

Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

Introducing randomness in strategies (2/2)

Several ways of **randomizing** $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$:

Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

Mixed strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

Introducing randomness in strategies (2/2)

Several ways of **randomizing** $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$:

Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

Mixed strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

General strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E))$$

Introducing randomness in strategies (2/2)

Several ways of **randomizing** $\sigma_\nabla : (V E)^* V_\nabla \rightarrow E$:

Behavioral strategies

$$\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E)$$

Mixed strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow E)$$

General strategies

$$\mathcal{D}(\sigma_\nabla : (V E)^* V_\nabla \rightarrow \mathcal{D}(E))$$

Kuhn's theorem¹⁵

All three classes are equivalent in games of *perfect recall*.

→ Requires access to infinite memory and infinite support for distributions.

¹⁵Aumann, "Mixed and Behavior Strategies in Infinite Extensive Games", 1964; Bertrand, Genest, and Gimbert, "Qualitative Determinacy and Decidability of Stochastic Games with Signals", 2017.

What about finite-memory strategies?

Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of memory states,
- ▷ m_{init} is the initial state,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the next-action function,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the update function.

What about finite-memory strategies?

Mealy machine $\mathcal{M} = (M, m_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:


- ▷ M is the set of memory states,
- ▷ m_{init} is the initial state,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow E$ is the next-action function,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow M$ is the update function.

Stochastic Mealy machine $\mathcal{M} = (M, \mu_{\text{init}}, \alpha_{\text{nxt}}, \alpha_{\text{up}})$:

- ▷ M is the set of memory states,
- ▷ $\mu_{\text{init}} \in \mathcal{D}(M)$ is the initial distribution,
- ▷ $\alpha_{\text{nxt}}: M \times V \rightarrow \mathcal{D}(E)$ is the next-action function,
- ▷ $\alpha_{\text{up}}: M \times E \rightarrow \mathcal{D}(M)$ is the update function.

⇒ **Three ways to add randomness: initialization, outputs, and updates.**

Taxonomy¹⁶ (1/2)



Classes **XYZ** with $X, Y, Z \in \{D, R\}$, where D stands for deterministic and R for random, and

- X characterizes the initialization,
- Y characterizes the next-action function,
- Z characterizes the update function.

¹⁶Main and Randour, "Different Strokes in Randomised Strategies: Revisiting Kuhn's Theorem Under Finite-Memory Assumptions", 2024.

Taxonomy (2/2)

This taxonomy holds from one-player deterministic games (no collapse) up to concurrent partial-information multi-player games (equivalences hold).

Taxonomy (2/2)


This taxonomy holds from one-player deterministic games (no collapse) up to concurrent partial-information multi-player games (equivalences hold).

↪ Collapses may arise for restricted classes of objectives (WiP).


Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.




From *home*, take the *train* or *bike* to reach *work*.




Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.




From *home*, take the *train* or *bike* to reach *work*.




Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.




From *home*, take the *train* or *bike* to reach *work*.




Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.




From *home*, take the *train* or *bike* to reach *work*.




Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.




From *home*, take the *train* or *bike* to reach *work*.




Multi-objectives MDPs (1/2)

We consider **two goals**:

- reaching work under 40 minutes with **high probability**;
- minimizing the **expectancy** of the time to reach work.



From *home*, take the *train* or *bike* to reach *work*.



Multi-objectives MDPs (2/2)

We are interested in the structure of this **payoff set**.

Our result¹⁷

For *good* payoff functions (\sim expectancy is well-defined),

- 1 the set of achievable payoffs coincide with the convex hull of *pure* payoffs;
- 2 we can approximate *any* strategy ε -closely by **mixing** a bounded number of *pure* strategies.

¹⁷ Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

Multi-objectives MDPs (2/2)

We are interested in the structure of this **payoff set**.

Our result¹⁷

For *good* payoff functions (\sim expectancy is well-defined),


- 1 the set of achievable payoffs coincide with the convex hull of *pure* payoffs;
- 2 we can approximate *any* strategy ε -closely by **mixing** a bounded number of *pure* strategies.

⇒ **RDD-randomization is sufficient in most multi-objective MDPs.**

¹⁷ Main and Randour, "Mixing Any Cocktail with Limited Ingredients: On the Structure of Payoff Sets in Multi-Objective MDPs and its Impact on Randomised Strategies", 2025.

Trading memory for randomness


Recall this generalized Büchi game asking to see *a* and *b* infinitely often:



We need (a two-state) memory to win it with *pure* strategies.

Trading memory for randomness

Recall this generalized Büchi game asking to see *a* and *b* infinitely often:




We need (a two-state) memory to win it with *pure* strategies.

But a (behavioral) randomized memoryless strategy suffices to win with probability one: playing v_1 and v_2 with non-zero probability ensures it.

Trading memory for randomness

Recall this generalized Büchi game asking to see *a* and *b* infinitely often:



We need (a two-state) memory to win it with *pure* strategies.

But a (behavioral) randomized memoryless strategy suffices to win with probability one: playing v_1 and v_2 with non-zero probability ensures it.

→ Memory can be traded for randomness for some classes of games/objectives.¹⁸

¹⁸Chatterjee, de Alfaro, and Henzinger, "Trading Memory for Randomness", 2004; Chatterjee, Randour, and Raskin, "Strategy synthesis for multi-dimensional quantitative objectives", 2014.

1 Controller synthesis

2 Memory

3 Randomness

4 Beyond Mealy machines

An incomplete story

Leitmotiv

Simpler strategies are better (for controller synthesis).

An incomplete story

Leitmotiv

Simpler strategies are better (for controller synthesis).

But what is simple?

An incomplete story

Leitmotiv

Simpler strategies are better (for controller synthesis).

But what is simple?

Usual answer: small memory, no randomness.

An incomplete story

Leitmotiv

Simpler strategies are better (for controller synthesis).


But what is simple?

Usual answer: small memory, no randomness.

↪ Let us question that.

Not all memoryless strategies are created equal


We want to reach v_3 .



Intuitively, the blue strategy seems simpler than the green one.

Not all memoryless strategies are created equal

We want to reach v_3 .




Intuitively, the blue strategy seems simpler than the green one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.

Not all memoryless strategies are created equal

We want to reach v_3 .




Intuitively, the blue strategy seems simpler than the green one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.
- ▷ The **representation of the next-action** function is mostly overlooked (basically a huge table).

Not all memoryless strategies are created equal

We want to reach v_3 .




Intuitively, the blue strategy seems simpler than the green one.

- ▷ Yet both are represented as a trivial Mealy machine with a **single memory state**.
- ▷ The **representation of the next-action** function is mostly overlooked (basically a huge table).

↪ **Memoryless strategies can already be too large to represent in practice!**


Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., **energy-Büchi** objective with $N \in \mathbb{N}$.



Memory is not always an issue


Multi-objectives games involving payoffs often require **exponential memory**. E.g., **energy-Büchi** objective with $N \in \mathbb{N}$.



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.

Memory is not always an issue


Multi-objectives games involving payoffs often require **exponential memory**. E.g., **energy-Büchi** objective with $N \in \mathbb{N}$.



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
- ↪ **Polynomial representation if we allow the use of counters.**

Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., **energy-Büchi** objective with $N \in \mathbb{N}$.




- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
- **Polynomial representation if we allow the use of counters.**

Hot take

We should explore novel notions of **simplicity**, and consider **alternative representations** of strategies/controllers.

Memory is not always an issue

Multi-objectives games involving payoffs often require **exponential memory**. E.g., **energy-Büchi** objective with $N \in \mathbb{N}$.



- ▷ We need a pseudo-polynomial Mealy machine because it lacks structure.
- **Polynomial representation if we allow the use of counters.**

Hot take

We should explore novel notions of **simplicity**, and consider *alternative representations* of strategies/controllers.

→ **We quickly survey a few ones in the next slides.**

Structurally-enriched Mealy machines

Idea:

- ▷ Augment Mealy machines with **data structures**: e.g., counters.¹⁹
 - ▷ Avoid “flattening” structural information about the strategy: more succinct representations, better understandability, and closer to actual controllers.
- ⇒ **Changes our way of thinking which strategies are complex or not.**

¹⁹ Blahoudek et al., “Qualitative Controller Synthesis for Consumption Markov Decision Processes”, 2020;
Ajdarów et al., “Taming Infinity One Chunk at a Time: Concisely Represented Strategies in One-Counter MDPs”, 2025.


Decision trees

- ▷ Structured state-space (e.g., $\subset \mathbb{Z}^n$) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

Decision trees

- ▷ Structured state-space (e.g., $\subset \mathbb{Z}^n$) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

Toy example: trying to reach the center $(0, 0)$ of a 2D-grid.



instead of

x	y	action
0	1	↓
0	2	↓
...	...	↓
-1	0	→
-1	1	→
...

Decision trees

- ▷ Structured state-space (e.g., $\subset \mathbb{Z}^n$) and action-space.
- ▷ Learn a (possibly approximative) decision tree from a given **memoryless** strategy.
- ▷ More understandable and compact than huge action tables.
- ▷ More complex tests may reduce size but hinder readability.

Works well in practice...²⁰

... starting from a given memoryless strategy.

²⁰Brazdil, Chatterjee, Chmelik, et al., "Counterexample Explanation by Learning Small Strategies in Markov Decision Processes", 2015; Brazdil, Chatterjee, Kretinsky, et al., "Strategy Representation by Decision Trees in Reactive Synthesis", 2018.

Other alternatives

■ Programmatic representations.

- ▷ Closer to realistic code, understandable.
- ▷ Strongly linked to the input format of the problem (e.g., PRISM code²¹), hard to generalize.

■ Models inspired by Turing machines.

- ▷ Powerful but hard to work with.
- ▷ Tentative notion of decision speed.²²

■ Neural networks.

- ▷ Prevalent in RL.
- ▷ Hard to understand and verify.
- ▷ Can be coupled with finite-state-machine abstractions.²³

²³Shabadi, Fijalkow, and Matricon, "Programmatic Reinforcement Learning: Navigating Gridworlds", 2025.

²³Gelderie, "Strategy machines: representation and complexity of strategies in infinite games", 2014.

²³Carr, Jansen, and Topcu, "Verifiable RNN-Based Policies for POMDPs Under Temporal Logic Constraints", 2020.

Wrap-up

Focus

Complexity of strategies in controller synthesis.

Wrap-up

Focus

Complexity of strategies in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

Wrap-up

Focus

Complexity of strategies in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ **Many questions are still open!**

Wrap-up

Focus

Complexity of strategies in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ Many questions are still open!

Strategy complexity \neq representation complexity.

Wrap-up

Focus

Complexity of strategies in controller synthesis.

Mealy machines are a powerful tool from a theoretical standpoint.

- ▷ High-level picture w.r.t. **memory** and **randomness**.

→ Many questions are still open!

Strategy complexity \neq representation complexity.

Take-home message

We need a **proper theory of complexity**, and a **toolbox of different representations**.

→ Ongoing project **ContolleRS**.

Thank you! Any question?